



Master's Degree Course in Computer Science

Master's Thesis

Lazy Declarative Heuristics in Answer Set Programming: Design and Evaluation of a clingo Propagator

Supervisor at the University of Calabria:
Prof. Carmine Dodaro

Candidate:
Pierpaolo Spadafora

Supervisors at the University of Potsdam:
Prof. Torsten Schaub
Dr. Javier Romero

Student ID 263722

Contents

Abstract	ii
1 Introduction	1
1.1 Context	1
1.2 Grounding Bottleneck and Dynamic Heuristics	1
1.3 Objective and Experimental Design	2
1.4 Contributions	3
1.5 Structure of the Thesis	4
2 Methodology	5
2.1 Conceptual Framework: Two Axes, Four Variants	5
2.2 System Architecture	7
2.3 The Native Backend	7
2.4 The Prolog Backend	10
2.5 Benchmark Problems and Encodings	11
2.6 Correctness Validation	12
2.7 Benchmark Infrastructure and Metrics	13
3 Results	14
3.1 Functional Validation	14
3.2 Experimental Setup and Dataset	14
3.3 Reduction of the Ground Heuristic Component	16
3.4 BSP: Runtime Behaviour	17
3.5 BSP: the Two Semantics Under the Microscope	20
3.6 BSP: Rules, Variables, and Memory	21
3.7 The Price of Laziness: Per-Decision Cost	22
3.8 Two Ground Simulations of Alpha: ga versus ga_weak	24
3.9 Direct and Auxiliary Forms	25
3.10 Effect of the Constraint Formulation	26
3.11 PUP: the Grounding Bottleneck at Scale	27
3.12 HRP: Semantics Without Aggregates	30
3.13 Native versus Prolog: Summary	31
3.14 Summary of Results	31
4 Discussion and Conclusions	33
4.1 Discussion of the Results	33
4.2 Limitations	34
4.3 Implications	35
4.4 Future Work	35
4.5 Conclusions	36

Abstract

Answer Set Programming (ASP) provides a declarative way to describe combinatorial problems while delegating the construction and exploration of the solution space to a grounding and solving system. In ground-and-solve systems such as clingo, the relevant part of a program is instantiated before search starts. This architecture is robust and mature, but it becomes expensive when heuristic directives generate a large number of ground instances, which happens precisely when the intended heuristic depends on the partial state of the computation, for example through an aggregate over the atoms that are currently true.

This thesis presents the design, implementation, and evaluation of a lazy heuristic propagator integrated into a modified version of clingo. The propagator keeps domain-specific heuristics in a compact, declarative form and evaluates them at solving time against the partial assignment, thereby avoiding the combinatorial materialization of large families of `#heuristic` directives. Two interchangeable evaluation backends are provided on top of the same propagation machinery: a *native* backend, in which heuristics are ground facts of the form `__heuristic(...)` interpreted by C++ runtime structures with incremental dynamic aggregates, and a *Prolog* backend, in which heuristics are rule strings of the form `heuristic("...")` evaluated by an in-process SWI-Prolog engine synchronized with the solver trail.

The evaluation is organized along two orthogonal axes. The first axis is the *approach*: ground-and-solve, where the heuristic is expanded by the grounder, versus lazy, where it is evaluated at runtime. The second axis is the *semantics of partial information*: a clingo-like reading, in which negation and aggregates refer to established falsity and to completely determined aggregate sources, versus an Alpha-like reading, in which `not a` is satisfied as soon as `a` is not assigned true and aggregates are computed over the atoms currently true in the partial assignment. Crossing the two axes yields the four variants `ga`, `gc`, `la`, and `lc` that are compared, together with a heuristic-free baseline, on three benchmark problems: Balanced Set Partitioning (BSP), the Partner Units Problem (PUP), and the House Reconfiguration Problem (HRP). A further technical contribution is the systematic *flattening* of Alpha's weight-priority pairs into single clingo weights, which is required because priorities rank heuristics globally in Alpha but only arbitrate modifications of the same atom in clingo.

The central experimental claim concerns memory. On BSP, the number of ground heuristic objects produced by the ground-and-solve variants grows as $\Theta(n^3)$, reaching 1,010,200 ground `#heuristic` directives at $n = 100$, whereas the lazy variants keep exactly two heuristic facts, independent of n . On PUP, the saving converts directly into reach: under a fixed 8 GiB budget, the lazy Alpha-like variant solves instances twice as large as the unguided baseline and one third larger than the best ground heuristic, whose unrolled directives exhaust memory first. On HRP, whose heuristics contain no aggregates, the comparison isolates the semantics of negation: the Alpha-like reading reproduces the intended guidance decision for decision, while the clingo-like reading proves either inert or actively misleading. Laziness does not eliminate the cost of the dynamic heuristic: it moves it from grounding-time space to solving-time work, a trade-off that the campaign quantifies through the per-decision metrics

of the propagator. The suite also includes an equivalence harness proving that all variants and both backends preserve the same answer sets, and guards that detect silent misconfiguration of the Prolog backend.

1 Introduction

1.1 Context

Answer Set Programming (ASP) is a declarative paradigm for knowledge representation and combinatorial problem solving. Instead of prescribing an algorithm step by step, the modeller describes the properties that admissible solutions must satisfy. The ASP system then computes the answer sets, that is, the stable models of the program, which correspond to the solutions of the encoded problem [1, 2].

Most modern ASP systems follow a *ground-and-solve* architecture. A grounder first rewrites the non-ground program, which may contain variables, into a propositional program. A solver based on CDCL-style techniques then explores assignments and computes stable models [2, 3]. clingo embodies this architecture by combining gringo, the grounder, and clasp, the solver. The separation is effective in many domains, but it also exposes a well-known limitation: grounding may become the bottleneck whenever the propositional instance grows much faster than the original declarative specification.

Domain-specific heuristics are one of the main tools for guiding search. In clingo they can be written as directives of the form:

```
1 #heuristic A : Body. [W@P, M]
```

Here A is the target atom, $Body$ is the applicability condition, W is the heuristic weight, P is a priority, and M is the modifier. The literature shows that domain-specific heuristics can be decisive on difficult industrial instances [4, 5]. At the same time, if a heuristic directive contains aggregates or a rich body, grounding the directive may produce a very large number of ground heuristic objects.

This thesis focuses precisely on that tension. The goal is to preserve a declarative form for heuristics, while moving part of their evaluation from grounding time to search time. The result is not a general lazy grounder for ASP. It is a specialized lazy mechanism for heuristics, implemented as a heuristic propagator inside clingo and instantiated in two interchangeable evaluation backends.

1.2 Grounding Bottleneck and Dynamic Heuristics

The grounding bottleneck appears when the propositional representation required by the solver becomes too large in time or memory. The phenomenon is especially visible in programs that combine numeric domains, aggregates, and conditions whose relevance depends on the evolving search state. Lazy-grounding research addresses this issue by interleaving grounding and solving, instantiating only the parts of the program that are needed during search [6].

Dynamic heuristics add a second layer of complexity. A useful heuristic choice may depend not only on the input instance, but also on the partial assignment currently built by the solver. A decision may become preferable because some atoms are already true, because other atoms are still undecided, or because an aggregate computed over true atoms has reached a favourable value. This idea is studied in

Alpha, a lazy-grounding system in which declarative heuristics can reason about partial states and dynamic aggregates [6, 7].

The reference papers make the source of the problem concrete. Consider the heuristic used throughout this thesis for Balanced Set Partitioning:

```
1  #heuristic b(X) : x(X), not c(X), S = #sum{Y : c(Y)}, W = ((n+1)*S)+X. [W,  
    ↪ true]
```

The aggregate value S appears *inside the weight*. Since the atoms $c(Y)$ are choice atoms rather than facts, the grounder cannot know the value of the sum at grounding time: it must enumerate every reachable value of S and materialize a distinct ground directive for every pair (X, S) . In a lazy-grounding system such as Alpha the same heuristic stays symbolic and the aggregate is simply re-evaluated over the atoms currently true. Reproducing this behaviour inside clingo, whose grounder is eager by design, is the engineering challenge addressed here.

clingo already offers a mature language for `#heuristic` directives and a Domain heuristic in the solver, but its ground-and-solve pipeline materializes the ground instances of such directives before search. When the heuristic component grows significantly, it becomes attractive to keep the heuristic in a compact representation and evaluate it at runtime.

1.3 Objective and Experimental Design

The objective of this thesis is to design and evaluate an extension of clingo that evaluates domain-specific heuristics lazily, and to compare it systematically against the ground-and-solve alternative. The comparison is organized along two orthogonal axes.

The first axis is the *approach*. In the *ground* approach (prefix `g`), the heuristic is expressed with native `#heuristic` directives and fully expanded by gringo. In the *lazy* approach (prefix `l`), the heuristic is kept as a small number of ground facts, read by a custom propagator during initialization and evaluated during search.

The second axis is the *semantics of partial information*, that is, the question of how a heuristic condition reads the current partial assignment. Under the *clingo-like* semantics (suffix `c`), negation and aggregates are evaluated as clingo would: `not a` is satisfied only when a is assigned false, and an aggregate refers to a completely determined set of source atoms. Under the *Alpha-like* semantics (suffix `a`), `not a` is satisfied as soon as a is not assigned true, and the aggregate is computed over the atoms currently true in the partial assignment, so its value evolves during search. Crossing the axes yields four variants per problem: `ga`, `gc`, `la`, and `lc`, evaluated against the heuristic-free baseline `gc_noheur`.

Two remarks are necessary to read this matrix correctly. First, everything runs on clingo: the variants labelled “Alpha” are *simulations* of Alpha’s semantics inside clingo, not runs of the Alpha system. In the ground variants, the simulation is necessarily approximate, because ground clingo cannot express a sum over the atoms currently true; the `ga` encodings therefore reconstruct the dynamic aggregate through a static unrolling combined with bound conditions, as explained in Section 2. Only

the lazy `1a` variant captures the dynamic aggregate exactly. Second, Alpha and clingo disagree on the meaning of weights and priorities: in Alpha the pair (W, P) ranks heuristics globally, with priority acting as a strict level discriminator between different atoms, whereas in clingo the priority `@P` only arbitrates modifications of the same atom and the global ranking is induced by the weight alone. Faithful translation therefore requires *flattening* the levels into the weight, $W_{flat} = W + P \cdot M$ for a sufficiently large level multiplier M , whenever the evaluation follows clingo’s reading of priorities.

The lazy approach is instantiated in two backends that share the same propagator interface. The *native* backend reads compact facts such as:

```

1  __heuristic(__target(b), x, __n_c, __bind(s, __sum(c)),
2          __weight(__add(__mul(s, B), self)), __priority(0),
3          __semantics(alpha), __modifier(true)) :- B = n + 1.

```

and evaluates them through C++ runtime templates with incremental aggregate states. The *Prolog* backend reads rule strings such as:

```

1  heuristic("heuristic(b(X), W, 0, true) :-
2      aggregate_all(sum(Y), holds(c(Y)), S), holds(heuristic_base(Base)),
3      holds(x(X)), target_available(b(X)), alpha_not(c(X)),
4      W is Base*S+X.").

```

and evaluates them with an in-process SWI-Prolog engine [8] kept synchronized with the solver trail. The two backends embody two different engineering trade-offs, compiled incremental evaluation versus a general logic-programming runtime, over the same declarative content.

The thesis addresses three questions. First, it asks how a dynamic heuristic can be represented in a compact ground form that a propagator can interpret, and how the Alpha-like and clingo-like readings of partial information can coexist in a single mechanism. Second, it studies how far the clingo semantics of weights, priorities, and modifiers can be reproduced through the `Clingo::Heuristic` interface, and what the weight–priority flattening costs in terms of encoding effort. Third, it evaluates the impact of the lazy representation on grounding size, memory, and search behaviour on three benchmark problems: Balanced Set Partitioning (BSP), the Partner Units Problem (PUP) [7], and the House Reconfiguration Problem (HRP) [6].

1.4 Contributions

The implemented system consists of two modified clingo builds sharing the same source layout. The build in `clingo-native` contains the native backend: shared types, a parser for `__heuristic` facts, incremental aggregate states, and the heuristic propagator, all under `libclingo`. The build in `clingo-prolog` contains the Prolog backend: the same propagator skeleton, a rule-string normalizer, an in-process SWI-Prolog evaluation backend, and an auxiliary clingo-based fallback evaluator. In both builds the propagator is registered automatically in `clingo_app.cc`, so the modified binaries use the mechanism without any external controller.

On the experimental side, the project provides encodings of BSP, PUP, and HRP in all four variants for both backends, under `test_folder/encodings-native` and `test_folder/encodings-prolog`; an equivalence harness, `tools/check_equivalence.py`, which verifies that heuristics preserve the answer sets and that the Prolog backend is actually active; and a benchmark pipeline built on `benchmark-tool` and `runlim`, with wrapper systems for the two binaries, a result parser that extracts solver and propagator statistics, and automatic graph generation.

1.5 Structure of the Thesis

Section 2 describes the methodology: the conceptual two-axis framework, the weight-priority flattening, the architecture of the two backends, the benchmark encodings with their adaptations, the correctness harness, and the measurement infrastructure. Section 3 presents the experimental results of the full campaign, executed on an HPC cluster under uniform resource limits: the complete analysis on BSP, the scaling study on PUP, the semantics study on HRP, and the cross-backend comparison with per-decision instrumentation. Section 4 discusses the findings, states the limitations explicitly, and outlines future work.

2 Methodology

2.1 Conceptual Framework: Two Axes, Four Variants

Every experiment in this thesis compares encodings of the same problem that differ only in how the domain-specific heuristic is represented and interpreted. The comparison is structured by two orthogonal axes.

The *approach* axis distinguishes where the heuristic is evaluated. In the ground-and-solve approach, the heuristic is written with native `#heuristic` directives; gringo expands them before search, and clasp’s Domain heuristic consumes the resulting ground objects. In the lazy approach, the heuristic remains a handful of ground facts; a custom propagator reads them at initialization and evaluates them at every decision point against the partial assignment.

The *semantics* axis distinguishes how a heuristic condition reads partial information. Under the clingo-like semantics, conditions are evaluated with respect to *established* truth values: a negative literal `not a` holds only if *a* has been assigned false, and an aggregate refers to the value it has once its source atoms are completely determined. Under the Alpha-like semantics, conditions are evaluated *optimistically* on the partial state: `not a` holds as soon as *a* is not assigned true, and an aggregate is computed over the atoms currently true, so its value changes as search proceeds. The Alpha reading is what makes heuristics genuinely *dynamic*: a directive such as “prefer the element whose insertion keeps the sums balanced” only makes sense if the sum is read off the evolving assignment.

Crossing the axes yields the four variants used throughout, summarized in Table 1, plus the baseline `gc_noheur`, which contains the domain logic without any heuristic.

	clingo-like semantics	Alpha-like semantics
Ground-and-solve	<code>gc</code> (native <code>#heuristic</code>)	<code>ga</code> (simulated, see §2.1)
Lazy (propagator)	<code>lc</code>	<code>la</code>

Table 1: The experimental matrix. Each cell exists for both evaluation backends (native and Prolog) and for each benchmark problem (BSP, PUP, HRP).

Two structural constraints shape the matrix. First, ground clingo cannot express the Alpha semantics exactly: there is no ground construct denoting “the sum of the atoms currently true”. The `ga` variants therefore *simulate* the Alpha reading, as described next, and only `la` captures it exactly. Second, the pair `gc/ga` is kept symmetric to the pair `lc/la`: within each approach, the two semantics differ exclusively in the treatment of negation and aggregates, so that any performance difference between suffixes can be attributed to the semantics and any difference between prefixes to the approach.

Simulating Alpha in Ground clingo

The `ga` encodings apply two systematic transformations to `gc`. The first transformation drops the negative literals on the partial state (for BSP, `not c(x)` and

`not b(X)`; for PUP, the `not not_assigned_...` conditions; for HRP, the guards such as `not fullCabinet(C)`). In Alpha these literals would also be satisfied on undecided atoms, so removing them is the closest ground approximation of the Alpha reading; the solver anyway never re-decides atoms that are already assigned.

The second transformation replaces exact aggregate bindings by *bounds over a static domain*. For BSP:

```

1 sum_value(0..tot).
2 #heuristic b(X) : x(X), sum_value(S), S <= #sum{Y : c(Y)},
3   W = ((n+1)*S)+X. [W, true]

```

Instead of binding S to the final value of the sum, the directive is unrolled over all candidate values `sum_value(S)` and guarded by the monotone condition $S \leq \#sum\{Y : c(Y)\}$. During search, as soon as the partial sum reaches a bound S , every directive with that bound becomes applicable; among the applicable ground directives for the same atom, clasp’s Domain heuristic retains the one with the maximum weight, which corresponds exactly to the *current* value of the growing aggregate. The max-weight selection thus reconstructs the dynamic aggregate at solving time. For aggregates whose contribution to the weight is decreasing rather than increasing (the free-capacity terms N_0, N_1, N_2 in PUP), the same trick is applied with the dual bound $V \geq \#max\{...\}$, so that the max-weight selection picks the smallest proven value. The price of the simulation is precisely the combinatorial explosion measured in Section 3: one ground directive per point of the product of the unrolled domains.

Weights, Priorities, and Flattening

Alpha and clingo interpret the annotation $[W@P]$ differently, and the difference is central to this work. In Alpha, heuristic annotations (W, P) form global levels: every applicable heuristic with priority P strictly precedes every heuristic with priority $P' < P$, regardless of weights, and the weight discriminates within a level. In clingo, the priority $@P$ only arbitrates between multiple heuristic modifications *of the same atom*; the global ranking among different atoms is induced by the resolved level value, that is, by the weight alone.

Encodings written for Alpha therefore cannot be transplanted into clingo verbatim. Whenever an ordering between different atom families is intended (zones before sensors in PUP; reuse before cabinet-filling before room-filling in HRP), the levels must be *flattened into the weight*:

$$W_{flat} = W + P \cdot M, \quad M > \max W + 1,$$

where M is a level multiplier strictly larger than any intra-level weight, so that no heuristic of a lower level can ever outrank a heuristic of a higher level. In the PUP encodings, M is the constant offset 100,000 added to the zone weights (the maximum sensor weight is bounded by 46,220 on the benchmark instances); in HRP, M is computed from the instance as `levelMul(LM) :- LM = MC + MR + 10` over the maximum cabinet and room identifiers.

The flattening rule interacts with the two axes as follows, and this interaction is enforced consistently across the suite. In the ground variants (`gc`, `ga`), clasp’s Domain heuristic is the consumer, so flattening is always required. In the lazy variants the consumer is the propagator, whose ranking behaviour is part of the semantics being modelled: under `__semantics(alpha)` the native backend ranks candidates globally by (priority, weight), faithfully reproducing Alpha’s levels, so `1a` encodings may use native priorities directly; under `__semantics(clingo)` the native backend deliberately reproduces clingo’s local-only reading of priorities, so `1c` encodings must flatten their weights exactly like `gc`. The Prolog backend, whose selection rule is a global (priority, weight) sort, expresses the same intended order through its priority field. The net effect is that all four variants of a problem implement the same intended decision order, each through the mechanism appropriate to its semantics.

A final convention concerns empty aggregates: `#max` over an empty set is 0 in Alpha but `#inf` in clingo. All encodings therefore use `#max{0; ...}` with an explicit neutral element, and both lazy backends implement the same convention (the native `MaxState` returns 0 when empty; the Prolog runtime defines `dyn_max` with a 0 default).

2.2 System Architecture

The project maintains two modified copies of clingo with the same source layout. The build in `clingo-native` implements the native backend; the build in `clingo-prolog` implements the Prolog backend and is configured with `CLINGO_USE_SWIPL=ON` so that the SWI-Prolog engine is linked in-process. In both builds the extension lives entirely under `libclingo` and is registered in `clingo_app.cc`:

```

1  HeuristicPropagator my_propagator;
2  ctl.register_propagator(my_propagator, true);

```

The class implements `Clingo::Heuristic`, the propagator interface that additionally exposes the `decide` callback. Both backends follow the same lifecycle: during `init` they scan the symbolic atoms for their trigger facts (`__heuristic(...)` in the native build, `heuristic(...)` in the Prolog build), build their runtime representation, and register watches on the solver literals they need to observe; during search, `propagate` and `undo` keep the internal state synchronized with the trail; at each decision point, `decide` either returns a signed literal for the best applicable heuristic target or returns 0, leaving the decision to the solver’s fallback heuristic. Runs are launched with `--heuristic=Domain`. Heuristics never change the answer sets: the propagator adds no constraints and only influences the order in which the solver explores the search space.

2.3 The Native Backend

Declarative Syntax

A native lazy heuristic is one ground fact per heuristic rule. The BSP heuristic, whose native directive would ground into $\Theta(n^3)$ objects, becomes:

```

1  __heuristic(__target(b), x, __n_c, __bind(s, __sum(c)),
2          __weight(__add(__mul(s, B), self), __priority(0),
3          __semantics(alpha), __modifier(true)) :- B = n + 1.

```

The arguments are, in order: the target predicate; an implicit positive body predicate (x , matched on the target tuple); a negative body predicate ($__n_c$ encodes $\text{not } c(x)$ on the target tuple); an aggregate binding that introduces the runtime variable s as the incrementally maintained sum over the source predicate c ; the weight expression, built from the operators `__add`, `__sub`, `__mul` over constants, bound variables, and the reserved symbol `self` (the first numeric argument of the target); the priority expression; the semantics selector; and the modifier. The ASP body `:- B = n + 1` is a small but important idiom: the fact is ground, and gringo evaluates the constant B and injects it as a plain number, so the propagator never needs to know program constants.

Richer bodies are expressed explicitly. When a positive body predicate does not share the target tuple, `__body(pred, __match(src, tgt), __bind_arg(var, idx))` specifies the argument correspondence and extracts values into runtime variables; `__bind_target_arg(var, idx)` names a target argument. Aggregates support per-target filtering with offsets, as in

```

1  __bind(m1, __max(num_sensors_on_unit, 0, __filter(1,1,-1)))

```

which, for a target `assigned_zone_unit(Z,U)`, maintains the maximum of the first argument of `num_sensors_on_unit(N,U')` restricted to source atoms with $U' = U - 1$. Since the aggregate machinery draws from a single source predicate, joins occurring inside Alpha-style aggregates are precompiled as auxiliary derived predicates in the encoding (for PUP, `sensor_assigned_zone/2` and `sensor_zone_on_unit/3`); these auxiliaries appear in no constraint, so they cannot change the answer sets.

Parser and Runtime Representation

The parser in `libclingo/src/heuristic_parser.cc` turns each `__heuristic` symbol into a `HeuristicRuleTemplate`: target predicate, positive body specifications, negative body predicates, variable bindings, modifier, semantics, and two arithmetic ASTs for weight and priority. Malformed syntax, duplicate variables, unknown operators, or negative indices are rejected eagerly with specific error messages, so that a misspelled heuristic fails fast instead of silently guiding nothing.

During `init`, the propagator builds an index from predicate names to the ground literals of their numeric tuples, and *materializes candidates*: for every ground target atom, and for every combination of matching positive body atoms, it creates a `RuntimeHeuristicCandidate` storing the target literal, the target tuple, the body literals, the values bound from body arguments, and the instantiated aggregate keys. Watches are registered on both polarities of the target and negative body literals, and on the positive body and aggregate source literals. Candidate materialization is proportional to the number of ground targets, but crucially it happens inside

the propagator’s memory as small POD structures, not as printed ground directives with their symbolic weights, and it does not multiply with the value domains of the aggregates.

Incremental Aggregates

Aggregate states implement a minimal interface with `add`, `remove`, and `result`. Sums and counts are plain accumulators; minima and maxima are ordered multisets, so that removal on backtracking is exact. Each source atom contributes its numeric value to the `RuntimeAggregateKey` obtained by instantiating the filter values; contributions are applied in `propagate` when a source literal becomes true and reverted in `undo`. Aggregates read under the clingo semantics are additionally *gated*: the propagator tracks the full set of source literals per runtime key and considers the aggregate usable only when every source is determined, which is the clingo-like reading of aggregate values. Under the Alpha semantics the current value is always usable.

Evaluation and Decision

A candidate is active when its target is still free, all positive body literals are true, every negative body literal is satisfied according to the template’s semantics,

```

1  static bool negative_literal_is_satisfied(HeuristicSemantics s,
    ↪  Clingo::TruthValue v) {
2      return s == HeuristicSemantics::Clingo ? v == Clingo::TruthValue::False
3          : v != Clingo::TruthValue::True;
4  }
```

and its variable environment can be built, which for clingo-semantics templates includes the aggregate gating above. Active effects are folded per target into a `TargetHeuristicState`. The modifiers follow clingo’s Domain heuristic: `level` sets the ranking value, `sign` the preferred polarity, `true/false` set both; `init` and `factor` are recognized but rejected, because they modify solver-internal scores that the `decide` interface cannot reach, and mapping them to `level` would be semantically misleading.

The per-target resolution and the global ranking implement the two readings of priority discussed in Section 2.1. For clingo-semantics effects, the local priority arbitrates between modifications of the same target, and the target enters the global ranking with its resolved weight only. For Alpha-semantics effects, the best effect per target is selected by (priority, weight), and the pair itself becomes the global rank of the target, reproducing Alpha’s levels. Ranked targets live in an ordered set:

```

1  std::set<DecisionRankKey, DecisionRankGreater> active_decision_ranks_;
```

ordered by decreasing priority, then decreasing weight, then a deterministic literal tie-break. `decide` pops stale entries lazily (targets that have been assigned, or whose state changed since ranking), applies the resolved sign to the best valid target, and otherwise falls back to the solver. `undo` uses `noexcept` refresh paths throughout, since exceptions during backtracking would violate the solver’s expectations.

2.4 The Prolog Backend

Rule Strings and Normalization

In the Prolog backend, a heuristic is a rule string whose head is `heuristic(Target, Weight, Priority,`

```
1 heuristic("heuristic(b(X), W, 0, true) :-  
2     aggregate_all(sum(Y), holds(c(Y)), S), holds(heuristic_base(Base)),  
3     holds(x(X)), target_available(b(X)), alpha_not(c(X)),  
4     W is Base*S+X.").
```

The body is genuine Prolog. The propagator collects all `heuristic/1` facts (the alias `prolog_heuristic/1` is accepted for compatibility), normalizes each string, and classifies its semantics syntactically: a rule containing `clingo_not(...)` is clingo-like, otherwise it is Alpha-like. To keep the synchronized state small, the propagator statically extracts the predicate signatures that actually occur in the rules (inside `holds`, `alpha_not`, `clingo_not`, `target_available`, and as rule targets) and synchronizes only atoms with those signatures.

The Runtime Program

At initialization the backend boots the SWI-Prolog engine in-process [8] and consults a generated runtime program that defines the evaluation vocabulary:

```
1 holds(A) :- true_atom(A).  
2 holds(A) :- static_atom(A), \+ true_atom(A).  
3 clingo_not(A) :- false_atom(A).  
4 alpha_not(A) :- \+ holds(A).  
5 target_available(A) :- target_atom(A), \+ true_atom(A), \+ false_atom(A).  
6 dyn_max(Goal, Template, Max) :-  
7     (aggregate_all(max(Template), Goal, R) -> Max = R ; Max = 0).
```

These few clauses are the entire semantic bridge. The solver trail is mirrored by the dynamic predicates `true_atom/1` and `false_atom/1`; a free atom is represented by absence from both. `alpha_not` and `clingo_not` implement exactly the two negation readings of Section 2.1; `aggregate_all` over `holds` computes dynamic aggregates on the atoms currently true, which is the Alpha aggregate semantics; `dyn_max/dyn_min` add the empty-set-is-zero convention; `target_available` guarantees that the rule only proposes targets that are still free.

Synchronization and Decision

Trail synchronization is incremental: `propagate` and `undo` update only the atoms mapped by the changed literals, each with a single combined retract-and-assert goal, since term parsing and calling dominate the synchronization cost. At each `decide`, the backend enumerates the solutions of `heuristic(T, W, P, M)`, discards candidates whose target is assigned or unmapped, and selects the best candidate by decreasing priority, then decreasing weight, then a deterministic tie-break; the modifier gives

the sign of the returned literal. The backend is enabled by the environment variable `LAZY_HEURISTIC_BACKEND=prolog`; with `LAZY_PROLOG_STATS=1` it prints a summary line with decision counts and cumulative timings of every phase (state synchronization, Prolog query, candidate scan, selection), which the benchmark parser ingests.

Without the environment variable, the same build falls back to an auxiliary evaluator that re-grounds a small clingo program per decision. This fallback accepts a different, ASP-style body syntax; a Prolog-style rule evaluated by it simply fails (SWI’s `unknown=fail` convention), so the run silently degenerates to no heuristic at all. This failure mode is invisible in exit codes and therefore dangerous for benchmarking; the guards that detect it are described in Section 2.6.

2.5 Benchmark Problems and Encodings

Balanced Set Partitioning (BSP)

BSP partitions $\{1, \dots, n\}$ into two sets `b/1` and `c/1` whose sums differ by at most one. The guessing part is a symmetric even loop, the check compares the two sums, and the heuristic prefers to place the next element into the set whose current sum is largest, thereby keeping the partition balanced greedily. Besides the four matrix variants and the baseline, BSP has three methodological extras: `ga_weak`, which drops the negative literals without the static unrolling and therefore changes only the activation moment; `la_aux`, which routes the lazy heuristic through an auxiliary predicate `h_b(X,S)` that reintroduces grounded aggregate values (and hence a grounding cost) before the propagator is even used; and `la_co`, which combines the lazy heuristic with a linear reformulation of the balance constraint, and is treated as a modelling experiment rather than a heuristic comparison.

Partner Units Problem (PUP)

PUP connects zones and sensors to communication units under adjacency and capacity constraints; the instances (the `double` and `doubleV` families) provide `comUnit/1` and `zone2sensor/2`. The encoding follows the Alpha evaluation encoding with an inductive, locality-based choice structure. The heuristic ranks sensor placements by a four-component weight (dynamic degree, forbidden placements, unit occupancy, direct connections) and zone placements by occupancy and free capacity, with zones globally preceding sensors.

Porting this encoding from Alpha to clingo required documented adaptations, applied uniformly to *all* variants so that comparisons remain fair: the inductive `assignable` rules are bounded by `comUnit(U#1)`, because gringo would otherwise materialize an unbounded chain that Alpha explores lazily; the existential capacity constraints are rewritten as monotone counting aggregates with linear grounding; a redundant distance-two blocking rule with cubic grounding is removed; counters are capped at the unit capacity (`sensorNumbers(0..2)`), without which the auxiliary predicate `num_forbidden_places_of_sensors` would explode in every variant and mask the effect under study. The zones-before-sensors ordering is flattened into the weight in `gc`, `ga`, and `lc` (offset +100,000), and expressed with a native priority in `la`, per the rule of Section 2.1.

House Reconfiguration Problem (HRP)

HRP extends the House Configuration Problem with legacy configurations: things must be stored in cabinets and cabinets placed in rooms owned by the things' owners, under capacity constraints refined by the thing-length/cabinet-size extension, while a legacy configuration can be partially reused or dismantled. The heuristic is a strict five-level policy: reuse legacy components first (with internal weights), then fill cabinets with long things, then with short things, then place cabinets into rooms, and finally a closing level that biases all remaining choice atoms to false, pruning the tail of the search. HRP is the negation-heavy benchmark: its heuristics contain no aggregates, so the `1a/1c` comparison isolates the effect of the negation semantics alone, while in BSP and PUP both mechanisms act together. In the native lazy encodings, the guards over the partial state are precompiled into same-tuple auxiliary predicates (for example `cabFull(C,T) :- cabinetDomain(C), thing(T), fullCabinet(C)`) so that negative conditions align with the target tuple; the five levels are native priorities in `1a` and are flattened with the instance-derived multiplier `levelMul` in `gc`, `ga`, and `1c`.

2.6 Correctness Validation

Heuristics must not change the answer sets: this is a falsifiable property, and the pipeline verifies it before measuring anything. The harness `tools/check_equivalence.py`, run by the sanity entry point ahead of every campaign, performs an *exhaustive equivalence* check on a small BSP instance: it enumerates all models with `clingo -n 0` for every variant and both backends, projects the models onto the solution predicates (`b/1`, `c/1`), and requires the projected collections to be identical across variants and across backends. Projection is necessary because the encodings expose different `#show` sets; comparing raw witnesses would produce false positives. For PUP, where exhaustive enumeration is intractable even on small instances, the harness checks agreement on SAT/UNSAT and reports the choice counts.

The harness also guards against the silent-fallback failure mode of the Prolog backend. Every lazy Prolog run must report `decide_calls > 0` in the propagator summary: a run that terminated successfully but never called the custom `decide` has degenerated to the no-heuristic baseline and is treated as an error by the harness (and flagged with an explicit warning column by the benchmark runner). As a cross-check, the harness verifies that native `1a` and Prolog `1a` perform the same number of choices on the control instance: if the SWI backend were inactive, the Prolog run would diverge from the native one. The benchmark wrapper for the Prolog system exports `LAZY_HEURISTIC_BACKEND=prolog` itself, so the guard protects against configuration drift rather than repairing it. The C++ unit tests complete the picture on the native side: they cover incremental aggregates, corner cases (empty domains, unsatisfiable instances), explicit body and target bindings, filtered aggregates, the local-versus-global reading of priorities in the two semantics, and the syntax validation of the parser.

2.7 Benchmark Infrastructure and Metrics

The campaigns are orchestrated with `benchmark-tool`, driving one run per (system, setting, instance) triple under `runlim` [9] for wall-clock and memory limits. The two systems are shell wrappers around the two modified binaries; both fix `--stats=2 --heuristic=Domain -n 1` and a numeric locale, and the Prolog wrapper additionally exports the backend activation and statistics variables. A custom result parser extracts, for every run: outcome and times (total, solving, and their difference as an estimate of the non-solving part), search statistics (choices, conflicts, restarts), grounding-size measures (ground `#heuristic` directives for the ground variants, lazy heuristic facts for the lazy ones, and rule counts from `clasp`’s statistics), the peak memory, and, for the Prolog backend, the propagator summary (decision calls, candidates seen per decision, cumulative synchronization, query, scan, and selection times).

Two measurement caveats are declared up front, because they define what the numbers can and cannot show. First, the memory metric is the peak resident set size of the whole solver process, end to end: for the lazy Prolog variants it includes the in-process SWI engine and its asserted database. The engine has a fixed initialization overhead independent of n , so on small instances the lazy variants may show a *larger* RSS than the ground ones; the lazy memory advantage is asymptotic, emerging where the $\Theta(n^3)$ growth of ground heuristics overtakes the constant overhead. Second, the ground-size comparison counts objects of different kinds: for ground variants it counts directives actually materialized by `gringo` (a faithful one-line-one-object count), whereas for lazy variants it counts template facts that expand at runtime into one candidate per ground target. The line count therefore *understates* the runtime work of the lazy backends by design; that work is measured separately by the per-decision metrics. Both caveats are taken up again in the discussion of the results.

3 Results

3.1 Functional Validation

The system was validated by keeping three levels separate. The first level concerns the ASP program itself: the answer sets, the SAT/UNSAT result, and the visible solutions must remain consistent across the compared encodings and across the two backends. The second level concerns the heuristic translation: targets, bodies, weights, priorities, modifiers, and aggregate bindings must be preserved when a native `#heuristic` directive is rewritten as a lazy fact, including the weight–priority flattening where the clingo reading of priorities applies. The third level concerns the effect on search, measured through choices, conflicts, solving time, memory, and the size of the ground heuristic component.

Heuristics do not change the semantics of the ASP program: they guide search, but they do not add or remove answer sets. The equivalence harness of Section 2 verifies this property exhaustively on BSP, by enumerating all models of every variant on both backends and comparing the collections after projection on `b/1` and `c/1`, and by SAT/UNSAT agreement on PUP, where exhaustive enumeration is intractable. The same harness enforces the anti-fallback guards for the Prolog backend (`decide_calls > 0` for every lazy Prolog run, and equal choice counts between native and Prolog `1a` on the control instance). On the native side, the C++ test suite isolates the correctness of the heuristic machinery from the benchmarks: it covers incremental aggregate states, empty and unsatisfiable corner cases, explicit body and target bindings, filtered aggregates, the local-versus-global interpretation of priorities under the two semantics, and the parser’s rejection of malformed syntax.

3.2 Experimental Setup and Dataset

The campaign reported in this section was executed on an HPC cluster under SLURM, with one job per (system, setting, instance) triple orchestrated by `benchmark-tool`. Every run was limited by `runlim` to 300 seconds of real time and 8 GiB of resident memory, on 4 dedicated cores. Both backends were compiled on the compute nodes against the same toolchain, and the Prolog backend embeds a modern SWI-Prolog (10.0.2) built in user space for the same environment. The dataset covers the three families with their canonical instance ranges: BSP with $n = 10$ to 120 in steps of 10; PUP with the `double- $N$` instances, $N = 20$ to 200 in steps of 20 communication units; HRP with the reconfiguration instances `house- $N$` , $N = 2$ to 20 persons in steps of 2. All five shared settings run on all families; the three methodological extras (`ga_weak`, `1a_aux`, `1a_co`) are BSP-specific. The grand total is 392 runs (192 for BSP, 100 each for PUP and HRP), none of which ended in error: every unsolved run is an explicit timeout (T) or memout (M).

The figures of this chapter follow the same separation. Every comparative plot shows the reference encoding and the four variants under study, `gc_noheur`, `gc`, `ga`, `1a`, and `1c`; the three exploratory settings are kept out of these plots on purpose, because each of them answers a different, narrower question, and appear only in the dedicated analyses of Sections 3.8, 3.9, and 3.10, in figures restricted to the variants that each comparison actually needs.

Two properties anchor the validity of the numbers. First, within each system the ground variants (**g***) execute identical binaries on identical ground programs, and clasp’s search is deterministic at fixed configuration: choice and conflict counts are therefore exact structural measures, not noisy samples, and they reproduce bit-for-bit the counts observed in the preliminary local campaign (for example, **gc** at $n = 100$ performs exactly 32,514 choices on both machines). Second, the anti-fallback guard confirmed `decide_calls` > 0 for every solved lazy Prolog run, so no lazy measurement silently degenerated to the baseline.

Tables 2–4 summarize the outcome of every cell of the experimental matrix. The columns report how many instances were solved and where the first failure occurs, distinguishing timeouts from memouts.

Variant	native		Prolog	
	solved	first failure	solved	first failure
gc_noheur	12/12	—	12/12	—
gc	11/12	T@120	11/12	T@120
ga	6/12	T@70	6/12	T@70
ga_weak	12/12	—	12/12	—
la	12/12	—	12/12	—
lc	12/12	—	12/12	—
la_aux	6/12	T@70; M@120	2/12	T@30; M@120
la_co	12/12	—	12/12	—

Table 2: BSP campaign overview ($n = 10..120$, 300 s / 8 GiB per run).

Variant	native		Prolog	
	solved	first failure	solved	first failure
gc_noheur	4/10	T@100; M@200	4/10	T@100; M@200
gc	3/10	T@80; M@140	3/10	T@80; M@140
ga	6/10	M from 140	6/10	M from 140
la	8/10	M from 180	8/10	M from 180
lc	4/10	T@80 (non-mon.); M from 180	2/10	T from 60; M from 180

Table 3: PUP campaign overview (**double- N** , $N = 20..200$). Note that **lc** is not monotone in N : the native run fails on **double-80** but solves **double-100**.

Variant	native		Prolog	
	solved	first failure	solved	first failure
gc_noheur	10/10	—	10/10	—
gc	10/10	—	10/10	—
ga	10/10	—	10/10	—
1a	10/10	—	10/10	—
1c	7/10	T from 16	7/10	T from 16

Table 4: HRP campaign overview (**house- N** , $N = 2..20$ persons).

Three headline facts are already visible at this granularity and are unpacked in the following subsections. The Alpha-like lazy variant **1a** is the most robust setting of the whole matrix: it is the only heuristic variant that never times out, on any family, on either backend. The ground Alpha simulation **ga**, which was designed as the faithful ground counterpart of **1a**, collapses on BSP precisely because of the static unrolling that makes it faithful. Finally, the clingo-like lazy variant **1c** exhibits two distinct failure modes, inertness on BSP and active misguidance on HRP, that the per-decision instrumentation makes directly measurable.

3.3 Reduction of the Ground Heuristic Component

The central claim of the thesis concerns *where* memory is spent, and the direct evidence is the number of ground heuristic objects as a function of n . The mechanism is readable in the encoding itself. The standard directive

```

1 #heuristic b(X) : x(X), not c(X), S = #sum{Y : c(Y)},
2   W = ((n+1)*S)+X. [W, true]

```

carries the aggregate value S as a variable inside the weight. Since the atoms $c(Y)$ are choice atoms, the grounder cannot know the sum at grounding time and must materialize one directive for every reachable pair (X, S) . With $X \in \{1, \dots, n\}$ and S ranging over the subset sums of $\{1, \dots, n\}$, which is every integer from 0 to $n(n+1)/2$ because 1 belongs to the domain, the count is

$$2 \cdot n \cdot \left(\frac{n(n+1)}{2} + 1 \right) = \Theta(n^3),$$

where the factor 2 accounts for the two symmetric rules for **b** and **c**. The formula is verifiable exactly on the measured data: for $n = 10$ it yields $2 \cdot 10 \cdot 56 = 1,120$, which coincides with the measured `ground_heuristics` value for **gc**. The explicit unrolling of **ga** materializes the same (X, S) product by construction. The lazy variants remove this enumeration at the root: the aggregate is evaluated at solving time on the current partial assignment, and the materialized heuristics remain constant, two facts, independent of n .




Figure 1: Comparison between native ground heuristic objects and lazy heuristic facts.

n	Native heuristic objects	Lazy facts	Ratio
10	1,120	2	560×
50	127,600	2	63,800×
100	1,010,200	2	505,100×
110	1,343,320	2	671,660×

Table 5: Growth of the native heuristic component compared with the lazy representation. The counts are properties of encoding and instance, independent of the machine.

This comparison must be read for what it is. For the ground variants the count is faithful: one line is one materialized object. For the lazy variants the two counted facts are *templates* that expand at runtime into one candidate per ground target (n candidates per rule for BSP); those candidates never pass through the grounder and therefore appear in no line count. The single-figure comparison “2 facts versus $\Theta(n^3)$ directives” is thus the proof of a saving in *grounding space*, not of a saving in total work: the work moved to runtime is real and is measured separately in Section 3.7.

3.4 BSP: Runtime Behaviour

Figure 2 reports the total time as measured by clingo, and Figure 3 its solving component. The base program dominates the picture for every same-encoding variant: at $n = 120$ the domain encoding grounds into 52.8 million rules, and the grounding component (about 160 seconds at $n = 120$) is essentially identical for `gc_noheur`, `la`, and `lc`, which share the encoding verbatim except for the heuristic representation.

What separates the variants is the solving component and, for the ground heuristics, the additional grounding of the directives themselves.




Figure 2: BSP, native backend: total clingo time. Curves are medians over solved instances; a curve ends where its variant stops solving within the limits.




Figure 3: BSP, native backend: solving time only. `la` is indistinguishable from zero at every size; `ga` leaves the chart after $n = 60$.

Variant	Total (s)	Grounding est. (s)	Solving (s)	Choices	Conflicts
<code>gc_noheur</code>	204.3	159.9	44.4	78,565	49,993
<code>gc</code>	<i>T</i>	(169.5)	(>130)	—	—
<code>ga</code>	<i>T</i>	(165.1)	(>134)	—	—
<code>ga_weak</code>	232.5	165.4	67.2	121,159	80,847
<code>la</code>	160.9	160.8	0.04	120	0
<code>lc</code>	207.8	160.7	47.1	78,565	49,993
<code>la_co</code>	0.005	0.005	0.00	117	0

Table 6: BSP at $n = 120$, native backend. *T* marks a timeout at 300 s; parenthesized values are the state of the run when it was killed. `la_aux` is discussed in Section 3.9.

Three observations follow from Table 6. First, `la` turns search into a formality, 120 choices and zero conflicts, so its total time is the grounding time of the base program: it is the fastest same-encoding variant at every size from $n = 60$ upwards and solves the whole range, whereas the ground directives push `gc` over the limit at $n = 120$ (the run was killed with at least 130 seconds already spent in solving on top of a grounding component inflated to 169 seconds).

Second, `ga` is the surprise of the full campaign, and a negative one. Its static unrolling was designed as the faithful ground simulation of the Alpha reading, and

Table 5 shows it pays the same $\Theta(n^3)$ grounding as `gc`; but unlike `gc`, whose negative guards deactivate directives as atoms become established, the unrolled directives of `ga` lack those guards by design, so the Domain heuristic is flooded with simultaneously applicable modifications. The search degenerates: at $n = 60$, `ga` needs 150 seconds against 8.5 of the baseline with 1.7 million choices, and from $n = 70$ it never finishes within the limit (at $n = 120$ it was killed at 238,397 choices and counting). The faithful simulation is thus *doubly* punished, once in the grounder and once in the solver: an instructive datum for anyone hoping to emulate dynamic heuristics in a ground-and-solve pipeline.

Third, Table 6 contains two rows that the figures of this section deliberately omit. `ga_weak`, the cheap approximation of the Alpha reading, stays close to the baseline and is compared against the faithful `ga` in Section 3.8; `1a_co`, which changes the problem model itself, is the subject of Section 3.10.

3.5 BSP: the Two Semantics Under the Microscope

Search statistics separate the two semantics sharply, and the full campaign adds an instrumentation-level proof of *why*. The `1a` variant makes exactly n choices with zero conflicts at every size: the Alpha reading makes the greedy balance heuristic total, one direct decision per element, which is the known optimal strategy for this problem. The `1c` variant is not merely “close to the baseline”: it is *search-identical* to it. At every size, its choice and conflict counts coincide with `gc_noheur` to the last unit (78,565 and 49,993 at $n = 120$).




Figure 4: BSP, native backend: choices and conflicts. `1a` lies on the horizontal axis; `1c` coincides with the baseline; `ga` leaves the chart after $n = 60$.

The Prolog instrumentation explains the coincidence. On BSP the `1c` rule guards the aggregate source with the clingo reading: the sum over `c/1` becomes usable only when every `c` atom is determined, and `clingo_not(c(X))` additionally requires established falsity. During search these conditions are essentially never satisfied at decision time: the backend reports an average of *zero* applicable candidates per decision (`avg_candidates_per_decide` = 0.0 over 78,565 decide calls at $n = 120$). Every decide call returns nothing, the solver falls back to its default heuristic, and the trajectory reproduces the baseline exactly. The clingo-like lazy variant on BSP

is therefore *inert by semantics*: it costs the full decide-loop machinery and buys nothing, which is precisely what the design predicts for a heuristic whose guards wait for information that only exists at the bottom of the search tree.

3.6 BSP: Rules, Variables, and Memory

The compression of the heuristic component is visible in the size of the propositional instance. At $n = 100$, **gc** is fifty times larger in solver variables than every other variant, because its heuristic directives materialize the negative body and the aggregate bound for each (X, S) pair. Notably, **ga** does *not* inflate the variable count (20,402, in line with the baseline): its unrolled directives reuse the same aggregate structures, and its cost shows up as grounding time and unguarded activations instead. The two ground variants thus pay the same combinatorial product in two different currencies.

Variant	Variables at $n = 100$
gc	1,030,606
ga	20,402
ga_weak	20,406
lc	20,300
la	20,300
gc_noheur	20,300

Table 7: Number of solver variables at $n = 100$, native backend.




Figure 5: BSP, native backend: peak resident set size.

The memory metric is the peak resident set size of the entire solver process, end to end: it covers grounding, CDCL search, and the custom propagator in a single value, and for the lazy Prolog variants it includes the in-process SWI-Prolog engine and its asserted database. At large n memory is dominated by the base program for every same-encoding variant: at $n = 120$, `gc_noheur`, `1a`, and `1c` sit at 6.74 GB, close to the 8 GiB limit, while `gc` reaches 7.6 GB before being killed by time. The heuristic saving is real but bounded by the base program’s own footprint; the variant that changes the footprint by orders of magnitude is `1a_co` (Section 3.10), because it changes the base program itself. On the Prolog side the SWI engine adds a fixed overhead of the order of ten megabytes, invisible at this scale and dominant only on the trivial `1a_co` runs. The RSS *confirms* the saving on the real consumption; the ground-object count of Table 5 *explains its cause*.

3.7 The Price of Laziness: Per-Decision Cost

Intellectual honesty requires stating that the lazy approach does not *eliminate* the cost of the combinatorial explosion: it *moves* it. What ground-and-solve pays once, in memory, at grounding time, the lazy backend pays repeatedly, in time, at every decision. The full campaign turns this from a declared principle into measured numbers, through the instrumentation of the Prolog backend (the native backend implements the same decision loop without per-phase timers, so the Prolog figures are an upper bound for it).

Family	Variant	decide calls	cand./decide	ms/decide	sync (ms)	overhead
BSP $n=120$	1a	120	121	0.50	3	+0.0%
BSP $n=120$	1c	78,565	0.0	0.24	4,349	+31%
PUP $N=160$	1a	197	5.3	3.79	8,384	+11%
HRP $N=20$	1a	106	1,248	8.46	34	$\times 4.4$
HRP $N=14$	1c	73,143	239	3.31	11,184	$\times 153$

Table 8: Prolog backend instrumentation on representative solved runs. *cand./decide* is the average number of candidates produced by the `heuristic/4` query per decision; *ms/decide* is the average query time; *sync* is the cumulative cost of mirroring the solver trail into the Prolog database; the last column compares the total time against the same variant on the native backend.

Table 8 decomposes the overhead along its two multiplicative factors: how often the heuristic is consulted (decide calls, driven by the search trajectory) and how much each consultation costs (driven by the number of candidates the rules generate and by the size of the synchronized state). The four corners of the space are all realized in the data. 1a on BSP is the benign corner: 120 consultations of half a millisecond each, a total overhead of well under one second on a 161-second run, so the Prolog and native totals coincide. 1c on BSP multiplies a *small* per-call cost by 78,565 inert consultations, adding 65 seconds (+31%) to a run in which the heuristic never contributes a decision, the worst possible bargain. 1a on HRP multiplies *few* consultations by a large per-call cost, since the Romero rules enumerate more than a thousand candidates per decision; on sub-second instances even that turns a 0.37-second native run into a 1.6-second Prolog run. 1c on HRP combines both factors, 73,143 consultations at 3.3 milliseconds each, and the query time alone (242 of 268 seconds) pushes it to the timeout that the native backend only reaches two sizes later. Synchronization, finally, scales with the trail activity and the number of watched atoms: negligible on BSP 1a (3 ms), it becomes the dominant lazy cost on PUP (8.4 seconds at $N=160$, against 0.75 seconds of query time).




Figure 6: BSP, Prolog backend: decide calls and per-decision query cost. The 1c curve pairs tens of thousands of calls with a low unit cost; 1a the opposite.




Figure 7: BSP: solving time, native (solid) versus Prolog (dashed) backend, lazy variants and ground reference.

The general law suggested by the table, and consistent with all 392 runs, is that the lazy overhead is the product $decide\ calls \times (query\ cost + amortized\ sync)$, where the first factor is governed by how well the heuristic guides (a good heuristic consults itself n times; a bad or inert one, once per baseline choice) and the second by how much of the program the rules can see. The practical reading is that *lazy heuristics are cheapest exactly when they work*: effective guidance keeps the decision count, and therefore the total overhead, small.

3.8 Two Ground Simulations of Alpha: ga versus ga_weak

The three subsections that follow analyze the exploratory settings. They are not alternative candidates in the main comparison: each one isolates a single design decision, and its figures include only the variants that decision touches.

The first decision is how faithfully the Alpha reading can be forced into a ground-and-solve pipeline. `ga` is the faithful simulation and pays for both halves of it: it drops the negative literals so that established truth does not deactivate the directives, *and* it statically unrolls the aggregate, replacing the equality binding of S with the monotone bound $S \leq \#sum\{Y : c(Y)\}$ over every candidate value, so that clasp’s max-weight selection reconstructs the current value of the growing sum at solving time (Section 2). `ga_weak` keeps only the first half: it treats the negation as Alpha does and leaves the aggregate untouched, which means the grounder evaluates the sum once, on the empty assignment, freezing $S = 0$ and degrading the heuristic

to a static ranking of the elements by X .




Figure 8: BSP, native backend, exploratory comparison: solving time (left) and solver choices (right) with both Alpha simulations alongside the main variants. `ga` degenerates at $n = 60$ (1.7 million choices, 142 seconds of solving) and times out from $n = 70$; `ga_weak` tracks the baseline over the whole range.

Figure 8 shows what each half costs. The unrolled `ga` explodes exactly where the unguarded activations flood the Domain heuristic: at $n = 60$ it needs 1.7 million choices and 142 seconds of solving against 0.8 seconds of solving (8.5 in total) of the baseline, and from $n = 70$ it never finishes, the double punishment already diagnosed in Section 3.4. `ga_weak` sits at the opposite end: structurally it is indistinguishable from the baseline (same rules, constraints, and variables, 20,406 versus 20,300 at $n = 100$, because nothing is unrolled), and behaviourally it stays close to it over the whole range (232.5 versus 204.3 seconds at $n = 120$, with somewhat noisier search: 121,159 choices against 78,565). It solves all twelve instances precisely because it renounces the dynamic part of the heuristic; the frozen $S = 0$ ranking neither guides nor obstructs.

Read together, the two simulations bracket the lazy propagator from both sides and confirm that the benefit measured for `1a` comes from the *dynamic* aggregate, not from the mere presence of weights or from the Alpha activation rule alone. The half of the Alpha reading that is cheap to ground (`ga_weak`) is behaviourally inert, and the half that reproduces the guidance (`ga`) is combinatorially unaffordable; only the runtime evaluation of the aggregate obtains the guidance without the unrolling, which is the design point the lazy propagator occupies.

3.9 Direct and Auxiliary Forms

The `1a_aux` variant routes the lazy heuristic through auxiliary predicates that materialize the heuristic body, including the aggregate value, before the propagator is used. This experiment separates two effects that are otherwise easy to confuse: reducing the number of ground heuristic directives, and avoiding the materialization of the heuristic body itself. `1a_aux` keeps only two `__heuristic` facts, but its auxiliaries `h_b(X,S)` and `h_c(X,S)` reintroduce the (X, S) product in the base program, and the full campaign shows the consequences from three angles at once. In the grounder,

the auxiliary atoms inflate the instance (145,600 solver atoms at $n = 50$ against 10,352 for **1a**). In the propagator, every (X, S) atom becomes a runtime candidate, so the per-decision scan grows with n^2 : the native backend times out from $n = 70$, and at $n = 120$ it exhausts the 8 GiB budget during initialization, before the first choice. On the Prolog side the same candidate flood must additionally cross the synchronization bridge, and the variant already fails at $n = 30$ (3.1 seconds of sync for 573 decisions at $n = 20$, against 3 *milliseconds* for **1a**). The lesson is clear: if the heuristic body is first materialized into a large auxiliary relation, the intended benefit of laziness is lost on every axis at once. The direct lazy form is the meaningful one.




Figure 9: BSP, exploratory comparison of the auxiliary form. Left: native solving time; **1a**_aux diverges from $n = 40$ and leaves the chart after $n = 60$, while the direct **1a** stays on the axis. Right: propagator decide calls on the Prolog backend; the auxiliary candidates multiply the consultations (573 at $n = 20$, against the n calls of the direct form), and the curve ends at $n = 20$, the last size the variant solves on that backend.

3.10 Effect of the Constraint Formulation

The **1a**_co variant changes not only the heuristic representation but also the BSP constraint itself, replacing the quadratic difference check over `sumB/sumC` with a single interval constraint over one sum, whose bounds are compile-time constants. The effect is dramatic and worth stating exactly: at $n = 120$ the propositional instance drops from 52,759,256 rules to 366, and the total time from 204.3 seconds (baseline) or 160.9 seconds (**1a**) to 0.005 seconds, with the same 117 guided choices and zero conflicts, on both backends (the Prolog run measures 0.096 seconds, the difference being the fixed engine startup). This measurement completes the diagnosis of Section 3.4: once the lazy heuristic has removed the search cost, the *entire* residual cost of BSP is the ground materialization of the balance check, and a formulation that avoids it removes the last Θ -large term. A comparison between **1a**_co and **1a** therefore measures the interaction between lazy heuristics and problem modelling, not the isolated effect of the heuristic representation, and it is reported as a modelling experiment: its role is to show what the combination “dynamic heuristic + propagation-friendly model” can reach when neither component pays a combinato-

rial ground price.




Figure 10: BSP, native backend, exploratory comparison of the constraint reformulation: grounding time (left) and propositional rules (right). `gc`, `la`, and `lc` overlap because they share the quadratic base encoding (52.8 million rules and about 160 seconds of grounding at $n = 120$); `la_co` lies on the horizontal axis of both plots (366 rules, milliseconds of grounding), the visual form of the two-orders-of-magnitude gap discussed in the text.

3.11 PUP: the Grounding Bottleneck at Scale

PUP is the family where the grounding bottleneck, not the search, sets the frontier, and where the lazy representation translates directly into solved instances. Figure 11 and Table 9 summarize the campaign.




Figure 11: PUP, native backend: total clingo time. Each curve ends at the last instance its variant solves.

Variant	max N solved	Total (s)	Grounding (s)	Peak RSS	first failure
gc_noheur	80	160.9	12.9	0.7 GB	T@100
gc	60	111.7	21.0	1.3 GB	T@80
ga	120	101.7	99.7	7.0 GB	M@140
1a	160	113.8	111.1	7.3 GB	M@180
1c	100	226.6	27.0	2.1 GB	T@80 (non-mon.)

Table 9: PUP frontier, native backend: largest `double-N` solved and the state of the run at that size.

The ordering of the frontiers tells the whole story. The baseline stalls at $N = 80$ by *time*: with no guidance, search explodes first (546,549 choices at $N = 80$, against 217 for `1a` at twice that size). The guided ground variants push further and hit the opposite wall: `ga` solves $N = 120$ in half the baseline’s time-at-frontier but dies of *memory* at $N = 140$, because its unrolled heuristic directives inflate the ground program to 7 GB (51 million rules at $N = 120$, against 12.6 million for the baseline at $N = 80$). `1a` dominates both: same search quality as `ga` (both semantics-Alpha, both near-zero conflicts: solving is 1.1 versus 2.0 seconds at $N = 120$), but with no heuristic grounding on top of the base program (3.4 GB and 47.7 seconds total at $N = 120$, that is, $2\times$ less memory and $2\times$ less time than `ga`). Its own frontier, $N = 160$, is set by the *base* program: 108.8 million rules and 7.3 GB of RSS for

the domain encoding alone, at which point no heuristic representation can help. In terms of reach, the lazy Alpha variant doubles the baseline frontier ($80 \rightarrow 160$) and adds one third over the best ground heuristic ($120 \rightarrow 160$), entirely by removing the heuristic’s grounding bill.




Figure 12: PUP, native backend: grounding time and peak memory. **ga** pays the heuristic unrolling (orange) on top of the base growth; **la** (green) pays the base program only, and meets the 8 GiB ceiling two sizes later.

The clingo-semantics cells complete the picture from below. Native **1c** oscillates around the baseline, solving $N = 100$ (which the baseline misses) but failing at $N = 80$ (which the baseline solves): as on BSP, its gated aggregates almost never produce candidates (the Prolog twin reports 0.04 candidates per decision), but here the rare decisions that do fire perturb the default trajectory unpredictably, in both directions. This is the expected behaviour of sporadic, semantically late guidance on a problem with high variance between neighbouring instances, and it is the reason the non-monotone frontier is reported explicitly in Table 3. On the Prolog backend the same inert consultations acquire a price (25,197 calls and 74 seconds of synchronization already at $N = 40$), and **1c** collapses to $N = 40$: the pathological corner of Section 3.7 reproduced at scale.




Figure 13: PUP: native-versus-Prolog solving comparison (left) and cumulative synchronization cost of the Prolog backend (right). For **1a** the synchronization, not the query, is the dominant lazy cost at scale.

For the well-behaved **1a**, the cross-backend comparison isolates the pure cost of the general-purpose engine: identical guidance (both backends solve exactly the same instances with near-identical search statistics), with the Prolog total growing from 113.8 to 126.0 seconds at $N = 160$ (+11%), of which 8.4 seconds are trail synchronization against 0.75 of query time. At PUP scale the bottleneck of the embedded engine is thus *keeping the state*, not answering the queries, which points directly at the batched-synchronization optimization discussed in future work.

3.12 HRP: Semantics Without Aggregates

HRP is the negation-heavy benchmark: its five-level Romero heuristic contains no aggregates, so the **1a/1c** comparison isolates the negation semantics alone. At the benchmark sizes the problem is easy for clasp, every variant except **1c** solves every instance in under half a second on the native backend, which makes HRP useless for frontier claims but ideal as a controlled experiment on guidance quality and per-decision cost.




Figure 14: HRP: solver choices on the native backend (left) and the native-versus-Prolog solving comparison (right). The only diverging curve is **1c**, in both plots.

The Alpha reading behaves exactly as designed: **1a** tracks **ga** decision for decision (86 versus 88 choices at $N = 14$, zero conflicts for both), confirming on a second problem class that the lazy propagator reproduces the intended Alpha guidance. The clingo reading, by contrast, is *actively toxic* on this heuristic: requiring established falsity on guards such as `not fullCabinet(C)` delays every level of the policy until the corresponding closure information exists, and the surviving low-level modifications steer the solver into a pathological region. Native **1c** needs 66,717 choices and 32,353 conflicts at $N = 14$ (almost eight hundred times the **1a** count) and times out from $N = 16$; the failure is identical on both backends, which certifies it as a property of the semantics, not of any evaluation engine. Where BSP showed the clingo reading as inert, HRP shows it as misleading: the two failure modes of “late” information are both on record.

The Prolog backend adds the cost dimension. The Romero rules enumerate every applicable (cabinet, thing) and (room, cabinet) pair, more than a thousand

candidates per decision at $N = 20$, so even the well-guided `1a` pays 8.5 milliseconds per consultation, and the misguided `1c` multiplies 3.3 milliseconds by 73,143 consultations: 242 of its 268 seconds at $N = 14$ are spent inside `aggregate_all`-free but candidate-rich Prolog queries, anticipating by two sizes the timeout that search alone would cause. The quantitative gap between the backends on `1c`, 1.75 versus 268 seconds at $N = 14$, is the largest in the whole campaign and is entirely attributable to the per-decision machinery, since the search trajectories coincide.

3.13 Native versus Prolog: Summary

Across the full campaign the two backends never disagree on *what* is solved except where the per-decision overhead itself is the binding constraint (`1c` on PUP, $N = 100 \rightarrow 40$; `1a_aux` on BSP, $n = 60 \rightarrow 20$; `1c` on HRP reaches the same frontier but two orders of magnitude slower). For every ground variant the two systems are byte-identical encodings on the same solver, and their times agree within measurement noise, which doubles as an end-to-end sanity check of the whole pipeline. For the lazy variants the native backend is uniformly at least as fast, with the gap governed by the law of Section 3.7; the Prolog backend remains the more expressive and the more observable one (its instrumentation produced every per-decision number in this chapter), which is the trade-off it was designed to embody.

3.14 Summary of Results

The full campaign supports seven conclusions. First, the lazy representation reduces the ground heuristic component from $\Theta(n^3)$ directives to two constant facts on BSP, and the saving is confirmed on real memory wherever the heuristic grounding is a significant share of the footprint, culminating in PUP where it converts directly into a doubled instance frontier ($80 \rightarrow 160$) under an 8 GiB budget. Second, the total runtime does not become constant, because the base program is still fully grounded: on BSP the grounding of the domain encoding dominates every same-encoding variant, and on PUP it is what ultimately stops even `1a`; the lazy mechanism isolates the heuristic from the growth of the instance, no more and no less. Third, the semantics axis is not a detail: the Alpha reading turns the BSP heuristic into a perfect greedy strategy (n choices, zero conflicts) and reproduces the intended Romero policy on HRP, while the clingo reading is inert on BSP (search-identical to the baseline, with zero applicable candidates per decision) and actively harmful on HRP; representation and semantics must never be conflated. Fourth, the faithful *ground* simulation of the Alpha reading (`ga`) is doubly punished at scale, in the grounder by the unrolling and in the solver by the unguarded activations, which makes the lazy propagator not merely cheaper but the only practical carrier of that semantics inside clingo. Fifth, laziness is a space–time trade-off whose runtime side is now measured, not merely declared: the overhead factorizes as decide calls times per-consultation cost, it ranges from unmeasurable (BSP `1a`) to dominant (HRP `1c`), and it is smallest exactly when the heuristic guides well. Sixth, the three exploratory controls delimit the approach from every side: `ga_weak` shows that the Alpha activation rule without the dynamic aggregate is behaviourally inert, `1a_aux` shows that materializing the

heuristic body forfeits the benefit entirely, and `1a_co` shows that pairing the lazy heuristic with a propagation-friendly model eliminates the residual ground cost and solves $n = 120$ in milliseconds. Seventh, the two backends agree on every scientific conclusion while differing exactly where their engineering differs, per-decision cost and observability, so the choice between them is a genuine trade-off rather than a correctness question.

4 Discussion and Conclusions

4.1 Discussion of the Results

The work shows that an important part of the cost of declarative dynamic heuristics can be moved from grounding to search. On BSP, native `#heuristic` directives grow as $\Theta(n^3)$ and exceed one million ground objects at $n = 100$; the lazy variants keep two facts, and the difference is visible in the propositional instance size, in the peak memory, and in the timeout profile. The mechanism generalizes, and the full campaign measures the generalization: on PUP, whose heuristics multiply four unrolled value domains per ground pair, the ground Alpha simulation exhausts an 8 GiB budget at `double-140` while the lazy variant carries the same guidance to `double-160` with half the memory and half the time at equal size, doubling the frontier of the unguided baseline. Where the heuristic grounding is the binding constraint, the lazy representation converts directly into solved instances; where the base program dominates, as at the top of both the BSP and the PUP ranges, it converts into an unchanged frontier reached at lower cost, and the residual wall belongs to the domain encoding, not to the heuristic.

This result should not be read as a replacement for general lazy grounding. clingo still grounds the base program; the propagator acts only on the heuristic component. Nor should it be read as a claim that laziness is free: the runtime work is real, it is paid at every decision, and the per-decision metrics exist precisely to keep that cost on the record. The honest formulation is a space–time trade-off with a crossover: below it, ground-and-solve remains competitive, because the explosion is manageable and the lazy machinery (in the Prolog case, an entire in-process logic-programming engine) carries a fixed overhead; beyond it, the constant-size heuristic component is the only representation that survives.

The second central finding is semantic, and it is a finding about *meaning*, not performance. The same declarative heuristic changes behaviour dramatically depending on how partial information is read. Under the Alpha-like reading, the BSP heuristic is a perfect greedy strategy (n choices, zero conflicts) and the HRP policy reproduces its intended level structure decision for decision. Under the clingo-like reading, the very same weights and bodies degenerate in two distinct ways that the campaign now documents separately: on BSP the guards never open at decision time, the propagator reports zero applicable candidates per call, and the search is bit-identical to the baseline (inertness); on HRP the surviving low-level modifications actively steer the solver into a region a thousand times larger, on both backends alike (misguidance). Neither reading is “wrong”: they answer different questions, and the four-variant matrix exists to keep those questions separate. A comparison that silently mixed the axes, for example by measuring `1a` against `gc` and attributing the difference to laziness, would be methodologically invalid; the thesis takes some care never to do so. The campaign adds a corollary about the ground side of the matrix: the faithful ground simulation of the Alpha reading, `ga`, is punished twice at scale, once in the grounder by the unrolling and once in the solver by the activations that its missing guards can no longer switch off, to the point that it is the only BSP heuristic variant that times out before the baseline does. Inside a ground-and-solve pipeline, the lazy propagator is thus not merely the cheaper carrier of the Alpha

semantics; at scale it is the only practical one.

The third finding concerns the translation layer between the two heuristic languages. Alpha’s weight–priority pairs are global levels; clingo’s priorities arbitrate on a single atom. Every encoding that expresses an ordering between different atom families must therefore flatten its levels into weights wherever clingo’s reading applies: in the ground variants, because clasp is the consumer, and in the lazy clingo variants, because the propagator deliberately reproduces clingo’s behaviour there. This is easy to state and easy to get wrong: an early revision of the lazy-clingo PUP and HRP encodings carried Alpha-style priorities that the clingo semantics does not rank globally, producing an unintended decision order that inverted the designed one (sensors before zones; cabinet-filling before reuse), and diverging from the Prolog twin encodings of the same variants. The final audit caught and corrected the discrepancy by applying the same flattening used in the ground encodings. The episode is instructive: the interaction between priority semantics and evaluation backend is exactly the kind of silent, non-crashing behavioural divergence that a benchmark pipeline cannot detect from exit codes, and it motivates the equivalence and wiring guards that the suite now includes.

A related engineering lesson comes from the silent-fallback failure mode of the Prolog backend. A single missing environment variable degrades every lazy Prolog run to the no-heuristic baseline without any error, since the fallback evaluator simply fails on the Prolog-style rule bodies. The pipeline defends itself with positive evidence of activity (`decide_calls > 0`), cross-backend agreement checks, and a wrapper that forces the variable; the general principle, that a heuristic mechanism must prove it ran rather than merely not fail, seems worth stating explicitly.

Finally, the auxiliary-form experiment delimits the approach from within. `la_aux` shows that laziness only pays if the *entire* dynamic part of the heuristic, aggregate included, stays out of the grounder: materializing the heuristic body into an auxiliary relation reintroduces the very product the mechanism was built to avoid.

4.2 Limitations

The current implementation has explicit technical limitations. The native backend indexes targets and bodies through numeric tuples only; non-numeric arguments are not supported. The arithmetic language covers addition, subtraction, and multiplication; the available dynamic aggregates are sum, count, min, and max, over a single source predicate, with joins delegated to precompiled auxiliary predicates in the encoding. The modifiers `init` and `factor` are recognized but deliberately rejected: in native clingo they modify solver-internal scores that the `decide` interface cannot reach, and approximating them as `level` would silently change their meaning. The propagator is registered sequentially and has no synchronization layer for parallel solving. The tie-break between targets of equal rank is deterministic but does not reproduce clingo’s internal scores.

On the evaluation side, the campaign now covers all three families on both backends with the corrected encodings, but its scope must be stated precisely. The per-decision instrumentation exists only in the Prolog backend: the native decision loop is not phase-timed, so cross-backend overhead is measured end to end and the

native per-phase profile is inferred, not observed. The HRP instances remain easy for clasp at every benchmarked size, so HRP supports conclusions about guidance quality and per-decision cost, not about frontiers. The PUP frontier claims rest on one instance per size (the `double` family): the non-monotone behaviour observed for `1c` is a reminder that neighbouring sizes differ in hardness, and per-size variance is not measured. The memory metric is end-to-end process RSS, which is the honest but blunt instrument discussed in Sections 2 and 3: it includes the Prolog engine for the lazy Prolog variants and cannot separate grounding memory from search memory. The auxiliary clingo evaluator of the Prolog build re-grounds a program per decision and is a functional reference, not a performance-comparable backend. Finally, the whitelist of static predicates and the inference of the program constant n in the Prolog backend are benchmark-specific conveniences that a production version should replace with a declarative interface.

4.3 Implications

The main implication is that declarative heuristics can be treated as an intermediate layer between modelling and solving. One does not have to choose between fully grounded `#heuristic` directives and procedural heuristics hard-coded in the solver: a compact declarative surface, evaluated at runtime with access to the partial assignment, is a workable third point in the design space, and it can host more than one evaluation semantics behind one syntax. The two backends demonstrate two implementation strategies for that layer, a compiled, incremental one and an embedded general-purpose engine, with the expected trade-off between per-decision efficiency and expressive freedom of the rule bodies.

From an engineering perspective, the work also shows that clingo’s propagator and heuristic interfaces are flexible enough to prototype non-trivial guidance mechanisms without touching the solver core: watched literals, incremental aggregate states, per-target modifier resolution, and a global decision ranking are all expressible against the public API of `Clingo::Heuristic`.

4.4 Future Work

Several extensions follow naturally. The mechanism should become configurable rather than always-on, and the expression language deserves integer division, modulo, and comparisons, together with a principled treatment of non-numeric symbols through a stable internal mapping. Aggregates with in-template joins would remove the need for precompiled auxiliary predicates. On the Prolog side, the synchronization protocol could move from per-literal goals to batched updates, and the backend could expose its per-decision statistics through clingo’s statistics interface instead of stderr parsing. A parallel-solving synchronization layer and a tie-break policy that cooperates with the solver’s own scores are natural solver-side improvements.

On the evaluation side, the PUP synchronization numbers make the batched-update optimization concrete and measurable: at `double-160` the trail mirroring costs eleven times the queries it serves, so a protocol that coalesces retract-assert pairs per propagation batch has a quantified target to beat. Beyond that, configu-

ration, scheduling, and packing domains are natural candidates for broadening the evaluation, precisely because their useful heuristics are dynamic; PUP would additionally benefit from multiple instances per size (the `doubleV` family) to put variance bars behind the frontier claims. A further direction is a verified converter from annotated `#heuristic` directives to lazy facts, with the flattening rule of Section 2 applied mechanically rather than by hand; the audit episode discussed above is a concrete argument for automating that step.

4.5 Conclusions

This thesis presented a lazy evaluation mechanism for declarative domain-specific heuristics in clingo, instantiated in two backends over a common propagator design: a native C++ backend interpreting compact `__heuristic` facts with incremental dynamic aggregates, and an in-process SWI-Prolog backend evaluating heuristic rules as logic programs against a synchronized image of the solver trail. The mechanism supports two explicit readings of partial information, the clingo-like and the Alpha-like semantics, and reproduces clingo’s weight, priority, and modifier behaviour where that reading applies, including the systematic flattening of Alpha’s global priority levels into clingo weights.

The full campaign on BSP, PUP, and HRP under uniform limits (300 seconds, 8 GiB) shows a reduction of the ground heuristic component from $\Theta(n^3)$ directives, more than one million at $n = 100$, to two constant facts, with the corresponding saving in real memory wherever the heuristic grounding matters; on PUP the saving doubles the solvable frontier of the baseline and outlives the best ground heuristic by one third. The campaign also quantifies the price: a per-decision runtime cost that ground-and-solve does not pay, which factorizes into consultation count times consultation cost, vanishes when the heuristic guides well, and dominates when it does not. The suite validates that all variants and both backends preserve the answer sets and guards against silent misconfiguration.

The final result is a working, documented, and measurable prototype together with a methodology: two orthogonal axes that keep representation and semantics separate, a flattening rule that makes Alpha-style encodings meaningful under clingo’s reading of priorities, and an experimental discipline that states what each number measures. It does not solve the grounding bottleneck in general, but it provides a concrete, verifiable path for bringing dynamic declarative heuristics into clingo without always paying the cost of their full ground materialization.

Bibliography

- [1] Martin Gebser et al. *Answer Set Solving in Practice*. Morgan and Claypool Publishers, 2012.
- [2] Martin Gebser et al. *Potassco Guide*. Second edition, version 2.2.0. University of Potsdam. 2019. URL: <https://potassco.org>.
- [3] Matthew W. Moskewicz et al. “Chaff: Engineering an Efficient SAT Solver”. In: *Proceedings of the 38th Design Automation Conference*. 2001.
- [4] Martin Gebser et al. “Domain-Specific Heuristics in Answer Set Programming”. In: *Proceedings of the Twenty-Seventh AAAI Conference on Artificial Intelligence*. 2013, pp. 350–356.
- [5] Carmine Dodaro et al. “Combining Answer Set Programming and Domain Heuristics for Solving Hard Industrial Problems”. In: *Theory and Practice of Logic Programming* (2016).
- [6] Richard Comploi-Taupe et al. “Domain-Specific Heuristics in Answer Set Programming: A Declarative Non-Monotonic Approach”. In: *Journal of Artificial Intelligence Research* 76 (2023), pp. 59–114.
- [7] Richard Comploi-Taupe, Gerhard Friedrich, and Tilman Niestroj. “Dynamic Aggregates in Expressive ASP Heuristics for Configuration Problems”. In: *Proceedings of the 25th International Workshop on Configuration (ConfWS 2023)*. Vol. 3509. CEUR Workshop Proceedings. CEUR-WS.org, 2023.
- [8] Jan Wielemaker et al. “SWI-Prolog”. In: *Theory and Practice of Logic Programming* 12.1-2 (2012), pp. 67–96.
- [9] Armin Biere. *runlim: a tool for sampling time and memory usage of a process*. <https://github.com/arminbiere/runlim>. Accessed 2026.

Acknowledgements

I would like to thank Prof. Carmine Dodaro for his guidance at the University of Calabria, and Prof. Torsten Schaub and Dr. Javier Romero for their supervision and support during my Erasmus period at the University of Potsdam. Their feedback helped shape both the modelling choices and the implementation decisions behind this work.

I am also grateful to the ASP and Potassco communities. Tools such as clingo make it possible to experiment with advanced research ideas while relying on a solid, well documented, and open foundation.

Finally, I thank my family and the people who stayed close to me throughout this thesis work, for their constant support and patience during the many long debugging sessions.